

Astrostatistics: Sat 10 Mar 2018

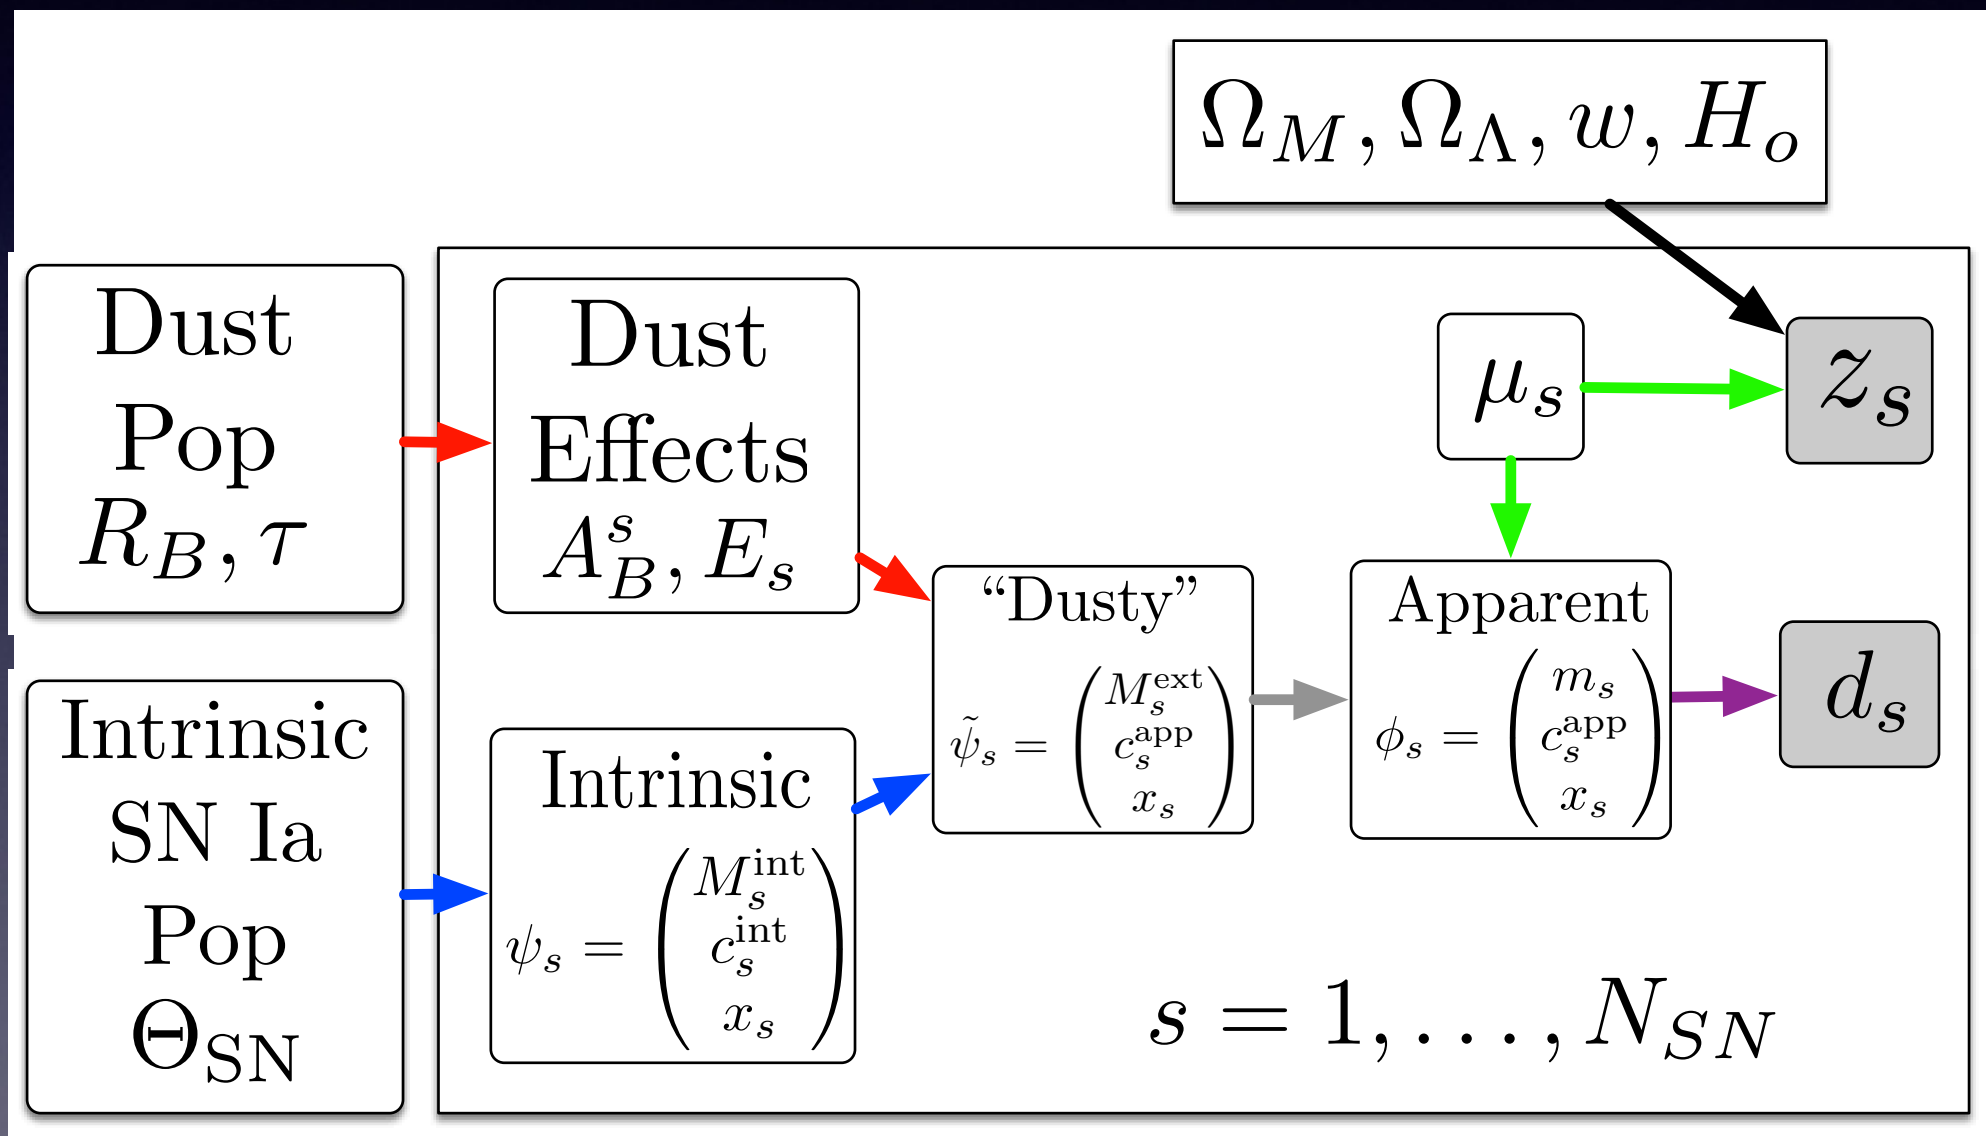
<https://github.com/CambridgeAstroStat/PartIII-Astrostatistics>

- Example Class 3: **Mon, Mar 12, 15:30~17:30pm MR14**
 - Quasar Time Delay Problem brings together Bayesian inference, MCMC, and Gaussian Processes
- Example Class 4 & Revision Class in Easter Term
- Email List & Course Questionnaires
- Last time: Gaussian Processes
- Today: Jeffrey's Prior & Start Probabilistic Graphical Models & Hierarchical Bayes
 - Gelman BDA, Chapter 5
 - Bishop, Pattern Recognition & Machine Learning, Chapter 8

Hierarchical Bayes, Huh?

What is it good for?

ABSOLUTELY EVERYTHING!



Mandel et al. 2017

Common Problems in Astronomy

- Want to learn about a population of objects from a finite sample of individuals, each measured with error
- Observed Data is actually a combination of uncertain astrophysical & instrumental & selection effects. Need to model them to infer the “intrinsic” properties of the object or population of objects (“deconvolve”)
- Research examples from the Class?

What is Hierarchical Bayes?

Simple Bayes: $\mathcal{D} | \theta \sim \text{Model}(\theta)$

Posterior (Bayes' Theorem): $P(\theta | \mathcal{D}) \propto P(\mathcal{D} | \theta)P(\theta)$

Hierarchical Bayes: θ_i : Parameter of Individual
 α, β : Hyperparameter of Population

$$\mathcal{D}_i | \theta_i \sim \text{Model}(\theta_i)$$

$$\theta_i | \alpha, \beta \sim \text{PopModel}(\alpha, \beta)$$

Joint Posterior:

$$P(\{\theta_i\}, \alpha, \beta | \{\mathcal{D}_i\}) \propto \left[\prod_{i=1}^N P(\mathcal{D}_i | \theta_i) P(\theta_i | \alpha, \beta) \right] P(\alpha, \beta)$$


Build up complexity by layering conditional probabilities

Probabilistic Graphical Models:

a visual way to understand complex statistical models

Simple Bayes:

$$\mathcal{D} | \theta \sim \text{Model}(\theta)$$




$$P(\theta | \mathcal{D}) \propto P(\mathcal{D} | \theta) P(\theta)$$

Hierarchical Bayes:

$$\mathcal{D}_i | \theta_i \sim \text{Model}(\theta_i)$$

$$\theta_i | \alpha, \beta \sim \text{PopModel}(\alpha, \beta)$$



$$P(\{\theta_i\}, \alpha, \beta | \{\mathcal{D}_i\}) \propto \left[\prod_{i=1}^N P(\mathcal{D}_i | \theta_i) P(\theta_i | \alpha, \beta) \right] P(\alpha, \beta)$$

Build up complexity by layering conditional probabilities

Probabilistic Graphical Models

Forward Model:

$$\theta_i | \alpha, \beta \sim \text{PopModel}(\alpha, \beta)$$

$$\mathcal{D}_i | \theta_i \sim \text{Model}(\theta_i)$$





Plate Notation:
(loop over
individuals in
sample)



$$P(\{\theta_i\}, \alpha, \beta | \{\mathcal{D}_i\}) \propto \left[\prod_{i=1}^N P(\mathcal{D}_i | \theta_i) P(\theta_i | \alpha, \beta) \right] P(\alpha, \beta)$$

Build up complexity by layering conditional probabilities

Advantages of Hierarchical Bayesian Models

- Common Problem in Astronomy: Infer properties of population from finite sample of individuals with noisy measurements
- Incorporate multiple sources of randomness & uncertainty as “latent variables” with distributions underlying the data
- Express structured probability models adapted to data-generating process (“forward model”)
- Bayesian: Full (non-gaussian) probability distribution = Global, coherent quantification of uncertainties
- Completely Explore & Marginalize Posterior trade-offs/degeneracies between parameters/hyperparameters

Simplest Hierarchical Bayesian / Multi-level Model: “Normal-Normal”

(Gelman BDA, Sec 5.4)

Level 1: Population Distribution of Latent Variables (Absolute Mags)

$$M_s \sim N(\tilde{M}, \tau^2)$$

Latent Variables Hyperparameters (Pop Mean & Variance)

Level 2: Measurement Error Process

$$\hat{M}_s | M_s \sim N(M_s, \sigma_s^2)$$

Measurements (Data) Heteroskedastic Meas. Error Variance (known)

(Draw PGM / DAG on Chalkboard) $s = 1 \dots N$

Hierarchical Bayesian “Normal-Normal” Model

Level 1: Population Distribution of
Latent Variables (Absolute Mags)

$$M_s \sim N(\tilde{M}, \tau^2)$$

Level 2 : Measurement Error
Process

$$\hat{M}_s | M_s \sim N(M_s, \sigma_s^2)$$

Joint Probability Density of Data, Latent Variables, Hyperparameters

$$P(\{\hat{M}_s\}, \{M_s\}, \tilde{M}, \tau^2)$$

Factor into Conditional and Marginal distributions based on Model

$$P(\{\hat{M}_s\}, \{M_s\}, \tilde{M}, \tau^2) = P(\{\hat{M}_s\}, \{M_s\} | \tilde{M}, \tau^2) \times P(\tilde{M}, \tau^2)$$

Hierarchical Bayesian “Normal-Normal” Model

Level 1: Population Distribution of Latent Variables (Absolute Mags)

$$M_s \sim N(\tilde{M}, \tau^2)$$

Level 2 : Measurement Error Process

$$\hat{M}_s | M_s \sim N(M_s, \sigma_s^2)$$

Joint Probability Density of Data, Latent Variables, Hyperparameters

$$P(\{\hat{M}_s\}, \{M_s\}, \tilde{M}, \tau^2) = P(\{\hat{M}_s\}, \{M_s\} | \tilde{M}, \tau^2) \times P(\tilde{M}, \tau^2)$$

$$P(\{\hat{M}_s\}, \{M_s\} | \tilde{M}, \tau^2) = P(\{\hat{M}_s\} | \{M_s\}; \tilde{M}, \tau^2) \times P(\{M_s\} | \tilde{M}, \tau^2)$$

(Likelihood)

$$P(\{\hat{M}_s\} | \{M_s\}; \tilde{M}, \tau^2) = P(\{\hat{M}_s\} | \{M_s\}) \quad (\text{Conditional Independence})$$

Hierarchical Bayesian “Normal-Normal” Model

Level 1: Population Distribution of Latent Variables (Absolute Mags)

$$M_s \sim N(\tilde{M}, \tau^2)$$

Level 2 : Measurement Error Process

$$\hat{M}_s | M_s \sim N(M_s, \sigma_s^2)$$

Likelihood function:

$$P(\{\hat{M}_s\} | \{M_s\}) = \prod_{s=1}^N P(\hat{M}_s | \{M_i\}) = \prod_{s=1}^N P(\hat{M}_s | M_s)$$

(Independent Measurements)

(Dist'n of Measurement s only depends on latent Mag s)

$$= \prod_{s=1}^N N(\hat{M}_s | M_s, \sigma_s^2)$$

Hierarchical Bayesian “Normal-Normal” Model

Level 1: Population Distribution of Latent Variables (Absolute Mags)

$$M_s \sim N(\tilde{M}, \tau^2)$$

Level 2 : Measurement Error Process

$$\hat{M}_s | M_s \sim N(M_s, \sigma_s^2)$$

Joint Probability Density of Data, Latent Variables, Hyperparameters

$$P(\{\hat{M}_s\}, \{M_s\}, \tilde{M}, \tau^2) = P(\{\hat{M}_s\}, \{M_s\} | \tilde{M}, \tau^2) \times P(\tilde{M}, \tau^2)$$

$$P(\{\hat{M}_s\}, \{M_s\} | \tilde{M}, \tau^2) = P(\{\hat{M}_s\} | \{M_s\}; \tilde{M}, \tau^2) \times P(\{M_s\} | \tilde{M}, \tau^2)$$

(Population Distribution)

$$P(\{\hat{M}_s\} | \{M_s\}; \tilde{M}, \tau^2) = P(\{\hat{M}_s\} | \{M_s\}) \quad \text{(Conditional Independence)}$$

Hierarchical Bayesian “Normal-Normal” Model

Level 1: Population Distribution of Latent Variables (Absolute Mags)

$$M_s \sim N(\tilde{M}, \tau^2)$$

Level 2 : Measurement Error Process

$$\hat{M}_s | M_s \sim N(M_s, \sigma_s^2)$$

Population Distribution (Parameterised Prior)

$$\begin{aligned} P(\{M_s\} | \tilde{M}, \tau^2) &= \prod_{s=1}^N P(M_s | \tilde{M}, \tau^2) \quad (\text{Each latent } s \text{ is independent, conditional on pop hyperparameters}) \\ &= \prod_{s=1}^N N(M_s | \tilde{M}, \tau^2) \end{aligned}$$

Hierarchical Bayesian “Normal-Normal” Model

Level 1: Population Distribution of Latent Variables (Absolute Mags)

$$M_s \sim N(\tilde{M}, \tau^2)$$

Measurement
Likelihood

Level 2 : Measurement Error Process

$$\hat{M}_s | M_s \sim N(M_s, \sigma_s^2)$$

Population Dist'n /
Prior

Joint Probability Density of ALL THE THINGS:
Data, Latent Variables, Hyperparameters

$$P(\{\hat{M}_s\}, \{M_s\}, \tilde{M}, \tau^2) = \prod_{s=1}^N [P(\hat{M}_s | M_s) P(M_s | \tilde{M}, \tau^2)] \times P(\tilde{M}, \tau^2)$$

Measurement
Likelihood

Population Dist'n /
Prior

Hyperprior

Putting the Bayesian in Hierarchical Bayesian

Joint Probability Density of ALL THE THINGS:
Data, Latent Variables, Hyperparameters

$$P(\underbrace{\{\hat{M}_s\}}_{\text{Data}}, \underbrace{\{M_s\}}_{\text{Unknowns}}, \tilde{M}, \tau^2) = \prod_{s=1}^N [P(\hat{M}_s | M_s) P(M_s | \tilde{M}, \tau^2)] \times P(\tilde{M}, \tau^2)$$

Joint Posterior of all unknowns given the data

$$P(\{M_s\}; \tilde{M}, \tau^2 | \{\hat{M}_s\}) = \frac{P(\{\hat{M}_s\}, \{M_s\}, \tilde{M}, \tau^2)}{P(\hat{M}_s)}$$

(Normalisation Constant)

Putting the Bayesian in Hierarchical Bayesian

Joint Posterior of all unknowns given the data

$$P(\{M_s\}; \tilde{M}, \tau^2 | \{\hat{M}_s\}) = \frac{P(\{\hat{M}_s\}, \{M_s\}, \tilde{M}, \tau^2)}{P(\hat{M}_s)}$$


 Unknowns Data (Ignorable Normalisation Constant)

(Posterior on N + 2 dimensional parameter space)

$$P(\{M_s\}; \tilde{M}, \tau^2 | \{\hat{M}_s\}) \propto$$

$$\left[\prod_{s=1}^N P(\hat{M}_s | M_s) P(M_s | \tilde{M}, \tau^2) \right] \times P(\tilde{M}, \tau^2)$$


 Measurement
Likelihood


 Population Dist'n /
Prior


 Hyperprior

Gibbs sampling & Hierarchical Bayes

Utilises Conditional Independence structure of PGM to
derive conditional posterior densities

$$P(\{M_s\}; \tilde{M}, \tau^2 | \{\hat{M}_s\}) \propto \left[\prod_{s=1}^N P(\hat{M}_s | M_s) P(M_s | \tilde{M}, \tau^2) \right] \times P(\tilde{M}, \tau^2)$$

1. Sample Latent Variables from Conditional: $s = 1 \dots N$

$$P(M_s | \tilde{M}, \tau^2; \{\hat{M}_s\}) \propto P(\hat{M}_s | M_s) P(M_s | \tilde{M}, \tau^2)$$

2. Sample Hyperparameters from Conditional:

$$P(\tilde{M}, \tau^2 | \{M_s\}; \{\hat{M}_s\}) = P(\tilde{M}, \tau^2 | \{M_s\}) \quad (\text{Conditional Independence})$$
$$\propto \left[\prod_{s=1}^N P(M_s | \tilde{M}, \tau^2) \right] \times P(\tilde{M}, \tau^2)$$